

RACE CONDITIONS

Software Security

Pedro Adão 2022/23

(with Ana Matos & Miguel Pupo Correia)

Motivation

Search Results (Refine Search)

There are **162** matching records.

Displaying matches **1** through **20**.

Search Parameters:

- **Keyword (text search):** race
- **Search Type:** Search Last 3 Years
- **Contains Software Flaws (CVE)**



1 [2](#) [3](#) [4](#) [5](#) [6](#) [7](#) [8](#) [9](#) [>](#) [>>](#)

CVE-2015-3247

Summary: Race condition in the worker_update_monitors_config function in SPICE 0.12.4 allows a remote authenticated guest user to cause a denial of service (heap-based memory corruption and QEMU-KVM crash) or possibly execute arbitrary code on the host via unspecified vectors.

Published: 9/8/2015 11:59:02 AM

CVSS Severity: 6.9 MEDIUM

CVE-2015-5189

Summary: Race condition in pcsd in PCS 0.9.139 and earlier uses a global variable to validate usernames, which allows remote authenticated users to gain privileges by sending a command that is checked for security after another user is authenticated.

Published: 9/3/2015 10:59:02 AM

CVSS Severity: 4.9 MEDIUM

CVE-2015-3212

Summary: Race condition in net/sctp/socket.c in the Linux kernel before 4.1.2 allows local users to cause a denial of service (list corruption and panic) via a rapid series of system calls related to sockets, as demonstrated by setsockopt calls.

Published: 8/31/2015 6:59:06 AM

CVSS Severity: 4.9 MEDIUM

Race conditions (I)

Attack: violation of an **assumption of atomicity**

- During a window of vulnerability or window of opportunity or window of inopportunity
- When 2 entities access concurrently the same object

Vulnerability: problem of concurrency / lack of proper synchronization

- between a target and malicious process(es); or
- between several processes/threads of the target

The attacker races to **break the assumption** during the window of vulnerability

Race conditions (II)

Example: service that produces unique sequential numbers

- Called by several threads concurrently
- Vulnerability allows returning the same number twice

```
int count = 0;    // shared
```

```
int getticket() {  
    count++;  
    -----  
    return count;  
}
```

Assumption of atomicity?
Window of vulnerability

Race conditions (III)

Sources of races

- Shared data: files and memory
- Preemptive routines (signal handlers)
- Multi-threaded programs

Mainly 3 kinds

- TOCTOU
- Temporary files
- Concurrency and reentrant functions

Example of running a race



ТОСТОУ

Quick review: symbolic links

Symbolic link, symlink or soft link

old concept, Posix, Unix (including Linux, macOS), Windows

Symlink is a special file that references another file or directory

full path of the other file

akin to a pointer or reference in programs

Create a link in a shell:

```
ln -s original_path link_path
```

Other operations equal to files: `rm`, `mv`

TOCTOU Time-of-check to time-of-use

aka TOCTTOU
also symlink attack

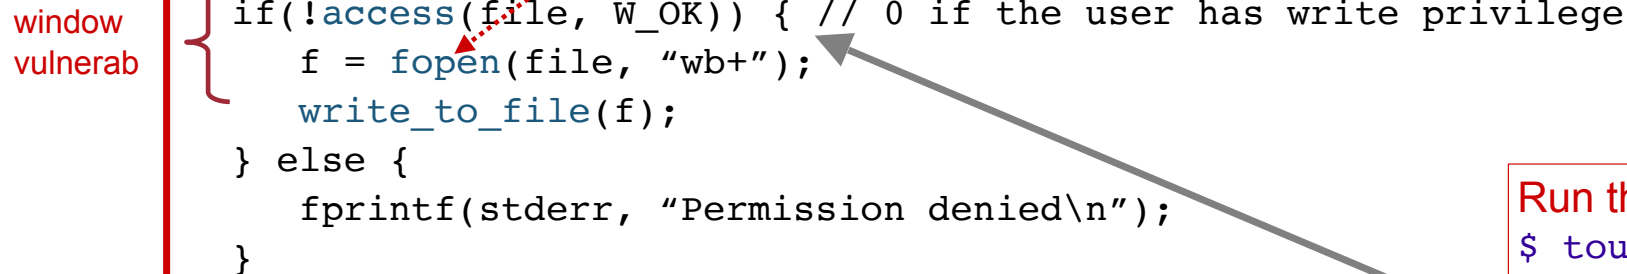
Typical case:

A program running *setuid* root is asked to write to a file owned by the user running the program

root can write to any file so the program has to check if the **actual** user has the right to write to the file

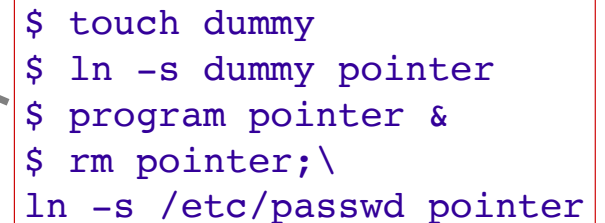
window vulnerab

```
{ if(!access(file, W_OK)) { // 0 if the user has write privilege
    f = fopen(file, "wb+");
    write_to_file(f);
} else {
    fprintf(stderr, "Permission denied\n");
}
```



Run this until success:

```
$ touch dummy
$ ln -s dummy pointer
$ program pointer &
$ rm pointer;\
ln -s /etc/passwd pointer
```



Program may test one file and open another

access

Designed for setuid programs (!)

Does **privilege check using the process' real UID** instead of the effective UID
i.e., checks if the user has access to the file, **even if the program is running with
effective UID <> real UID**

However it is vulnerable to race conditions / TOCTOU

The file it checks can be altered before it is used

Why? Because it is identified by a path, so the object in the end can change

Should never be used

A real example (I)

Broken `passwd` command (old, SunOS, HP/UX)
takes as parameter the password file (!) then:

- window of vulnerability {
1. Open password file, retrieve user's entry, close file
 2. Create and open temporary file (`ptmp`) in the same directory as password file
 3. Open password file again, copy content to `ptmp` and update modified info
 4. Close both files, rename `ptmp` to be the password file

A real example (II)

Attack script (interleaved with desired passwd execution):

```
mkdir evil
cp /etc/passwd evil/passwd
echo "hacker::0:0:::/bin/bash" >> evil/passwd
ln -s /etc link
passwd link/passwd
```

passwd step 1: open `link/passwd`, get user entry, close file

passwd step 2: create `ptmp`


```
rm link ; ln -s evil link
```

passwd step 3: copy `link/passwd` to `ptmp`, update modified info

```
rm link ; ln -s /etc link
```

passwd step 4: close files; move `ptmp` to `link/passwd` (i.e. `/etc/passwd`)

`/etc/passwd` has a new user `hacker` with superuser privileges



Putting it to work

Running `passwd` and exploiting the window of vulnerability seems hard...

Solutions:

- Make a script that tries as many times as needed until the attack works
- Delay the program somehow
- Insert some random delays between operations

Easier if the file system is distributed (e.g., NFS)

stat family

The culprits behind many TOCTOUs

```
int stat(const char *pathname, struct stat *buf);
```

pathname – file to be checked

buf – structure to be filled with data about the file

lstat – the same but data returned is about the link, if pathname is a link

These functions give information about the file

Owner, owning group, number of hardlinks to the file

Type of file (regular, dir, char device, block device, named pipe, symbolic link, socket)

They test the file/link, but when it is used it may be another one

Problem is atomicity, not order (I)

setuid program has to open a file but does not want to be tricked into opening a symbolic link

idea: to use `lstat()` instead of `access()`

Vulnerable code in Kerberos 4 lib used in login daemon:

```
if (lstat(file, &statb) > 0) goto out;  
if (!(statb.st_more & S_ISREG) || ...) goto out;  
if ((fd=open(file, O_RDWR|O_SYNC, 0)) < 0) goto out;
```

is it a regular file (not a link)?

If file is not a link it is safe to open...

but what if it changes after `lstat()` and before `open()` ?

Problem is atomicity, not order (II)

Still vulnerable if done in the opposite order (**open** first)

```
fd = open(fname, O_RDONLY);  
if (fd == -1) perror("open");  
if (lstat(fname, &stbl) != 0) die("file not there");  
if (!S_ISREG(stbl.st_mode)) die("it's a symlink");
```

Attacker creates link, then substitutes it by some file after the **open()** and before the **lstat()**

The file that is opened and processed is the one that is linked

Some other file is checked and understood not to be a link

The problem is atomicity, not order

Preventing file race conditions (I)

Most file races have to do with the resolution of pathnames

When a call with a pathname is done (`open`, `access`, `stat`, `lstat`,...), the pathname is resolved until the inode is found

If two calls are made one after the other
the path can lead to different *inodes*

Solution: avoid the two sequential resolutions by avoiding using **filenames** inside the program

Preventing file race conditions (II)

Correct example – with file descriptors

```
fd=open("/tmp/bob", O_RDWR);  
fstat(fd, &sb);
```

If someone unlinks and re-links `/tmp/bob` between the two calls, `fd` would still point to the **same** inode

a file descriptor is an integer, which is an index for an entry in a kernel array data structure containing the details of open files

Unsafe: `access`, `stat`, `lstat`, `chmod`, `chown`

Safe: `fstat`, `fchmod`, `fchown`

TEMPORARY FILES

Temporary Files

Have the same problems as others plus those derived from usually being in a shared dir

`/tmp, /var/tmp`

Typical attack:

1. Privileged program checks that **there is no file X in /tmp**
2. Attacker races to **create a link called X** to some file, say `/etc/passwd`
3. **Privileged program** attempts to create X and **opens the attacker's file** doing something undesirable that its privileges allow...

mktemp, tmpnam,...

`mktemp()` creates unique, currently unused, filename from template

```
strcpy(temp, "/tmp/tmpxxx");  
if (!mktemp(temp)) die("mktemp");  
fd = open(temp, O_CREAT | O_RDWR, 0700);  
...
```

After `mktemp()` and before `open()` an attacker can link `temp` to some file
`open()` would then not create but open an existing file

`tmpnam` and `tempnam` are similar

NAME [top](#)

tempnam - create a name for a temporary file

SYNOPSIS [top](#)

```
#include <stdio.h>
```

```
char *tempnam(const char *dir, const char *pfx);
```

Feature Test Macro Requirements for glibc (see [feature_test_macros\(7\)](#)):

tempnam():

Since glibc 2.19:

 _DEFAULT_SOURCE

Glibc 2.19 and earlier:

 _BSD_SOURCE || _SVID_SOURCE

DESCRIPTION [top](#)

Never use this function. Use [mkstemp\(3\)](#) or [tmpfile\(3\)](#) instead.

mkstemp, tmpfile,...

`mkstemp` - much safer than `mktemp ()` since it

1. **atomically checks for uniqueness,**
2. **creates and**
3. **opens with rw privileges**

`mkdtemp ()` is similar but creates a directory

Temporary files (cont)

Problematic solutions:

Random file names - often the attacker can try the race many times...

Locks - many implementations are enforced by convention, not mandatory

Solutions:

Use long random numbers (e.g. 64 bits)

set *umask* appropriately (e.g. 0006)

open with *fopen*

Safe calls seen before

CONCURRENCY AND REENTRANT FUNCTIONS

Concurrency

In the previous cases, *concurrency is created by the attacker*, with malicious intention

In many cases in which there is normal concurrency of operations on objects, operations may have to be executed **atomically**, i.e., without interruption

Solution: mutual exclusion mechanisms

locks, mutexes, semaphores, transactions, etc.

Problems:

1. **Starvation:** a thread is never scheduled for execution (problem of fairness)
2. **Deadlock:** threads inter-block themselves
3. **Race conditions...**

Reentrancy (I)

A function is **reentrant**

Definition: A function is reentrant if it works correctly even if its thread is interrupted by another thread that calls the same function

i.e., if several instances of the function can be executed in parallel in the same address space

Example of a non-reentrant function supposed to give unique tickets

```
int count=0;    //global var
int getticket() {
    count++;
    return count; } atomic?
}
```

Reentrancy (II)

Similar notions

Thread-safety – reentrancy in relation to multi-threading

Async-signal-safety – reentrancy in relation to signals

Sufficient conditions for function being reentrant:

(i.e., if these are satisfied, then function is reentrant)

Cannot use: static variables, global variables, other shared resources like libraries (i.e., uses only local non-static variables and function parameters)

Can only call reentrant functions

Necessary conditions for function being reentrant:

The sufficient conditions don't need all to be satisfied for the function to be reentrant

Using static/global variables or other shared resources is ok if they are only read

Signal handlers

Signals are a Unix/POSIX mechanism used to indicate asynchronous events to a process

Can be treated by a signal handler (a function) or ignored or blocked

Signal handlers have to be **asynchronous-safe** (or **signal-safe**):

Have to run safely and correctly even if interrupted by asynchronous events (i.e., by a signal)

Otherwise they may be vulnerable, and often are

Signal vulnerability example

```
char *user;
int cleanup(int sig) {
    printf("Cleaning up\n");
    free(user);
    exit(1);
}

int main(int argc, char **argv) {
    signal(SIGTERM, cleanup);
    signal(SIGINT, cleanup);
    ...
    process_file(fd);
    free(user);
    close(fd);
    printf("bye!\n");
    return 0;
}
```

```
int process_file(int fd) {
    char buffer[1024];
    ...read from file into buffer...
    user=malloc(strlen(buffer)+1);
    strcpy(user, buffer);
    ...
    return 0;
}
```

What if there is a signal...

1. after `free(user)`: there is a double free (corrupts heap)
2. during `strcpy`: data would be copied into a free heap block (mem corruption, arbitrary code exec)
3. anywhere: if SIGINT then SIGTERM fast (or vice-versa) → double free

Signals

Some solutions

- Make signal handler simple; ideally only set a flag

- Use system calls safe for signal handlers

- Block signals

 - inside signal handlers and

 - during non-atomic operations in the program

Java servlets

Usually

Called by several threads concurrently (typically)

Each servlet is only one object in memory...

...so class attributes are shared resources (like global vars in C)

Bad solution: implementing SingleThreadModel interface; bad performance

Servlet that can give two users the same count:

```
public class Counter extends HttpServlet {  
    int count = 0;    // shared!!  
    public void doGet(..., HttpServletResponse out) throws...{  
        PrintWriter p = out.getWriter();  
        count++;  
        -----  
        p.println(count + " hits so far!");  
    }  
}
```


Solutions

```
public class Counter extends HttpServlet {
    int count = 0;    //shared!!
    public synchronized void doGet(..., HttpServletResponse out)
        throws...
    {
        PrintWriter p = out.getWriter();
        count++;
        p.println(count + " hits so far!");
    }
}
```

not efficient

(if method was longer...)

```
public class Counter extends HttpServlet {
    int count = 0;    //shared!!
    public void doGet(..., HttpServletResponse out) throws... {
        int my_count;
        PrintWriter p = out.getWriter();
        synchronized(this) { my_count = ++count; }
        p.println(my_count + " hits so far!");
    }
}
```

Prevention

Identify resources that are used by several threads

Typically global variables, global data structures

Use proper mutual exclusion mechanisms

Sometimes the mechanisms are used but

e.g. a lock is released too early

Example: `sys_uselib` in the Linux kernel used a semaphore but released it before a call that should also be in the critical section → a race could be used to get root privileges

Summary

Race conditions are **violations** of assumptions of **atomicity** during a ***window of vulnerability***

Usual objectives are *privilege escalation* and attacks against web applications

Common vulnerabilities:

- TOCTOU – File access

- Temporary files

- Non-reentrant functions – signals, servlets