

Dynamic protection

Software Security

Pedro Adão 2022/23

(with Ana Matos & Miguel Pupo Correia)

Dynamic protection

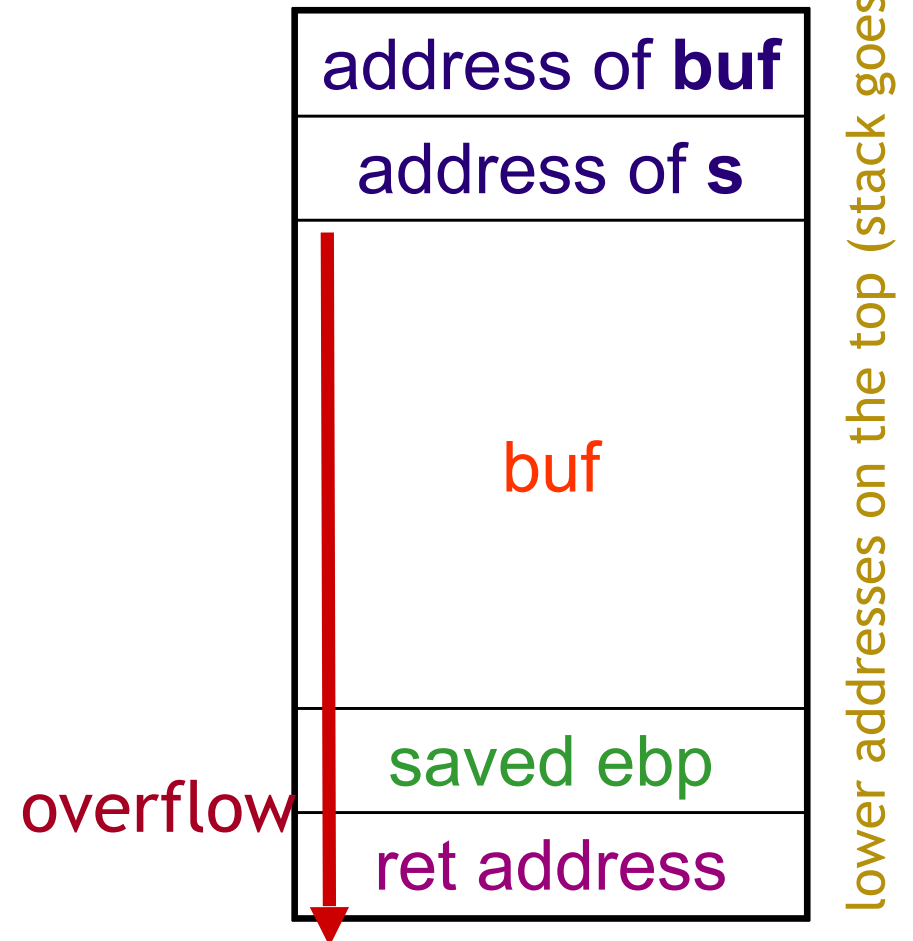
- The idea is to block attacks that may exploit existing vulnerabilities
- We consider mostly the case of memory corruption attacks
 - They are one of the most pervasive classes of attacks: buffer overflows, format strings
 - Good idea to assume that they will not completely disappear...
 - so mechanisms to minimize their effects when they do happen are important
 - The topic is vast, there are many others

Detection with canaries

Canaries / Stack cookies

```
void test(char *s) {  
    char buf[10];  
    strcpy(buf, s);  
    ...  
}
```

- Stack smashing:

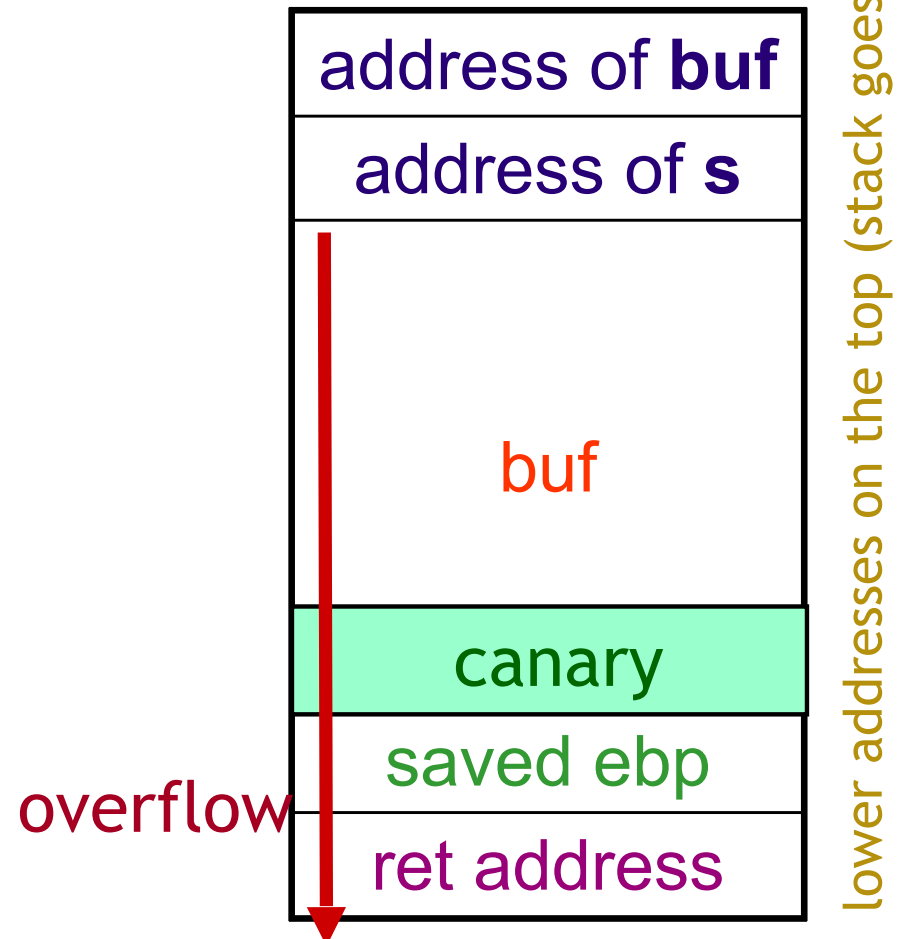


Canaries / Stack cookies

```
void test(char *s) {  
    push canary;  
    char buf[10];  
    strcpy(buf, s);  
    ...  
    if (canary is changed)  
        {log; exit;;}  
}
```

- Idea: canaries in the mines
- Compile time technique
- Canary = 32bit random value
- Appeared in StackGuard; /GS flag in Visual Studio 7.0 and above
- Implemented by the compiler
 - e.g., -fstack-protector/
-fno-stack-protector in gcc

- Stack smashing:



Canaries / Stack cookies (cont)

- Detects some stack smashing attacks
 - Stack smashing attacks that **overwrite the return address** (saved EIP)
 - to jump to injected code or ret-to-libc
 - **Off-by-one errors** that overwrite the saved EBP
- Do **not** detect BO attacks that **modify local variables** (they are *above* the canary)
 - Function-pointer clobbering, Data-pointer modification, Exception-handler hijacking (Windows Structured Exception Handler - SEH)
- Detects, but possibly too late, BO attacks against the function arguments (are *below* the ret address)


above = smaller address (in the stack); below = higher address (in the stack)

Canaries / Stack cookies (cont)

- Vulnerable to denial-of-service attacks that crash the app on purpose
- Compile time and some code (e.g. libraries, 3rd party code) cannot be recompiled; but most code surely can


Local variables protection

- Canaries do not protect local vars from overflows
- Solution: reorder the stack layout
 - All char buffers are put below all other local variables
 - Ex: ProPolice




Local variables protection

- Canaries do not protect local vars from overflows
- Solution: reorder the stack layout
 - All char buffers are put below all other local variables
 - Ex: ProPolice




Function arguments protection

- BO may overwrite function's arguments in the stack
 - Canary detects the attack after these arguments were used in the function
- 2 solutions
 - Some compilers solve this automatically: they put the args in CPU registers, preventing this attack (but there aren't many registers)
 - Otherwise: in the beginning of the function, copy the args to the top of the stack
- Currently Windows' /GS flag does that (either regs or top of stack)



Function arguments protection

- BO may overwrite function's arguments in the stack
 - Canary detects the attack after these arguments were used in the function
- 2 solutions
 - Some compilers solve this automatically: they put the args in CPU registers, preventing this attack (but there aren't many registers)
 - Otherwise: in the beginning of the function, copy the args to the top of the stack
- Currently Windows' /GS flag does that (either regs or top of stack)



Non-executable stack and heap

Non-executable stack and heap

- Many buffer overflow attacks involve injecting shell code in the stack/heap
 - Simple protection: mark these memory pages as non-executable (NX)
 - Old idea but x86 CPUs did not support it until ~2004
 - Several names:
 - Intel - Execute Disable (XD-bit)
 - AMD - Enhanced Virus Protection
 - Microsoft - Data Execution Prevention (DEP)
 - others call it W^X (write or execute, but not both)
 - Activated by the compiler (e.g., `-z noexecstack` in gcc); usually on by default

Non-executable stack and heap (cont)

- Limitations of NX
 - Doesn't protect from ret-to-libc/return-oriented programming (ROP) attacks
 - There may be lib functions that can be called to turn-off NX
 - ex: in a Windows program a call to `NtSetInformationProcess` disables NX (to allow apps to load DLLs incompatible with NX)
 - not too critical as attacker would have to first exploit a vulnerability to call that function
 - Some apps can be incompatible with it
 - Apps that perform high-performance image manipulation/rendering sometimes create binary code in runtime
 - Some interpreted languages compile scripts into binary code

Randomization and obfuscation

Address space layout randomization

- Idea is to randomize the addresses where code and data are placed in runtime
 - This is not what normally happens: the memory layout tends to be the same for every execution
 - What is randomized are *not* the *physical addresses* by shuffling pages around the RAM (which always happens anyway), but the logic/linear addresses, i.e., the organization of the virtual memory of a process
 - Doesn't prevent exploitation, makes it unreliable/harder
 - Started with PAX and Exec Shield projects for Linux and in OpenBSD, now in many others including MS Windows

ASLR (cont)

- Elements that can be randomized
 - Code - addresses where apps and dynamic libraries are loaded
 - Stack - starting address of the stack of each thread
 - Heap - base address of the heap
- Not all bits are randomized to reduce fragmentation
 - Example after two reboots of Windows:
Kernel32 loaded at 77AC0000 LoadLibrary at 77B01E7D
Kernel32 loaded at 77160000 LoadLibrary at 771A1E7D
 - Program that prints var address re-executed 4 times:

Mac OS X (32 bits)	Linux Ubuntu (32 bits)
end: 0xbff <u>c</u> 8b7e	end: 0xbfd <u>e</u> 1c7f
end: 0xbff <u>c</u> 5b7e	end: 0xbfc <u>b</u> a <u>e</u> 1f
end: 0xbff <u>1</u> 0b7e	end: 0xbfe <u>f</u> e0df
end: 0xbff <u>f</u> bb7e	end: 0xbfd <u>a</u> 18cf

ASLR (cont)

- What/when is randomized? (Windows)
 - User apps and some DLLs randomized whenever loaded
 - Shared DLLs (e.g., kernel32.dll) randomized only once per reboot
 - Old apps and DLLs not randomized
 - Because randomization requires that the executables include relocation information, which was not generated by old linkers
- ASLR is effective against most BO attacks
 - Stack smashing (*w/shellcode injection*), ret-to-libc
 - **Not for those against local variables**

ASLR (cont)

- Limitations
 - Anything with a static location is candidate to counter ASLR (e.g., apps or DLLs without relocation information)
 - Works only at reboot for many DLLs, therefore a local attack can first discover the memory layout, then do the main attack
 - Brute force attacks possible if target code restarts on failure
 - especially if the random space is small (e.g., 256 options in Windows)
 - => number of restarts should be limited

Instruction set randomization

- Code injection would be almost impossible if each computer had its own random instruction set (!)
- Randomized instruction set
 - Legitimate code is scrambled (e.g., XORed with a random value), then unscrambled for execution
 - Malicious code is not scrambled, but it is unscrambled, becoming (hopefully) impossible to execute
 - Scramble best at load time; at compile time, key would be fixed
 - Unscramble has to be done by a virtual machine or debugger (high overhead, not practical) or by the CPU (not available yet)

Instruction set randomization (cont)

- More practical case: SQL instruction randomization to avoid SQL injection
 - Application adds a key to each SQL command and operator
 - e.g. if key is 333 then:
 - \$query = “SELECT * FROM orders WHERE id=” . \$code; becomes:
 - \$query = “SELECT333 * FROM333 orders WHERE333 id=333” . \$code;
 - Between the web server and the DBMS there is a proxy that checks if all instructions follow the format (i.e., if they include 333). If not an error is generated
 - Example of query modified by attack:

OR does not follow format **OR333**
 - SELECT333 * FROM333 orders WHERE333 id=3331 **OR 1=1**
- Same idea might be applied with other languages to avoid **code injection**

Function pointer obfuscation

- Long-lived function pointers (typ. global vars) are often the target of memory corruption exploits
 - Provide a direct method for seizing control of program execution
 - E.g., heap overflow /format string attack that can only write a 32-bit value
- Pointer obfuscation mitigates this problem
 - The idea is to XOR the pointer with a random secret cookie
 - Keep it protected while not needed, unprotect when needed
 - Windows: EncodePointer(), DecodePointer(), EncodeSystemPointer(), DecodeSystemPointer()
- Not a complete solution, but useful with ASLR and NX
- Others have proposed checking if a function pointer points to the code segment

Integrity verification

SEH protection

- Windows Structured Exception Handling (SEH)
 - When an exception is generated, Windows examines a linked list of EXCEPTION_REGISTRATION structures **in the stack**
 - These structures include **pointers to the handlers** → can be overrun and an exception raised to force a jump to that address

SEH protection (cont)

- /SafeSEH compilation flag mitigates the problem
 - A *module* (DLL, .exe) can have an *exception table* that lists the valid EXCEPTION_REGISTRATION structures for the module
 - The validity of exception handler is checked before it's called
 - If the *module* has an *exception table* and the EXCEPTION_REGISTRATION structure is not registered there, the check fails;
 - otherwise the check is successful (if it matches or no table)
- Limitations:
 - if the *module* has no *exception table*...

Array bound checking

- BOs caused by lack of array bound checking, so doing the check solves the problem
 - Already done in languages like Java
 - Currently supported in C compilers like gcc (flag *Warray-bounds*)
 - However, it's easy to check `a[3]`, but not `*(a+3)` or `a[i]`

gcc gives warning


```
char str[10];
int i;
for (i=0; i<10; i++) {
    str[i] = 'a';
}
str[10]=' \0';
```

gcc doesn't give warning

```
char str[10];
int i;
for (i=0; i<10; i++) {
    str[i] = 'a';
}
str[i]=' \0';
```

Detection through interception

- *libsafe* - library with wrappers for problematic libc functions
- When a function is called, the wrapper first checks if the buffer is not being overrun
- How? Use *frame pointer* (EBP) as an upper-bound



Control-flow integrity

- Stack Shield has a global ret stack:
 - Whenever a function is called, the return address is stored in a Global Ret Stack (besides the normal stack)
 - Before the function returns, check if the return address in the normal stack and the Global Ret Stack match
 - Global Ret Stack has by default only 256 entries, so allows a maximum depth of 256 function calls
 - Code has to be added at compile time

Control-flow integrity (cont)

- Control flow: ifs, whiles, jumps, calls,...
 - Return from functions is just one case and global ret stack a partial solution
- Generic control-flow integrity checking
 - A static analysis tool annotates the program with data about valid destinations of jumps. Ex:

```
void (*f)() = ...;    /* function pointer */  
...  
f(); //annotated with valid destinations
```
 - Before jumping, check if the destination is valid

Filtering

Web application firewalls

- WAFs are application-level firewalls for webapps



Web application firewalls



- **Interposition** can be obtained in 4 ways
 - **Bridge**: WAF installed as a transparent bridge (switch)
 - **Router**: network reconfigured to direct traffic through the WAF
 - **Reverse Proxy**: represents the web server for the clients.
 - **Embedded**: WAF is installed as a web server plug-in
 - Ex: the most popular open source WAF, ModSecurity, works in modes 3 or 4 (with the Apache web server)

Sw or Hw box

WAFs (cont)

- HTTPS / encrypted traffic; 3 solutions:
 - HTTPS channel ends at the WAF (instead of the server)
 - Server provides decryption key to the WAF
 - Not an issue if the WAF is embedded
- Filtering criteria
 - Negative model: WAF has description of bad input
 - Positive model: WAF has description of good input
 - 2nd is best in theory (fail-safe defaults); 1st may be easier in practice in many cases (less false positives)

WAFs (cont)

- Models are sets of signatures/rules; can be:
 - Generic for certain attacks
 - Ex. authentication brute forcing, XSS, SQLI,...
 - Specific for a webapp
 - Ex. specific: block access to Perl CGIs present but not used by the app, block access to admin pages
- Origin of the signatures/rules:
 - Come with the WAF (not specific for the web app)
 - best WAFs are those with best models (ModSecurity)
 - Created manually (using the WAF language)
 - Machine learning

WAFs (cont)

- Reaction to the detection of an attack
 - Block HTTP request
 - Break TCP connection
 - Block the source IP address and break all its TCP connections
 - Block the webapp session
 - Block the webapp user

Summary

- BO detection with canaries
- Non-executable stack and heap
- Randomization and obfuscation
- Integrity verification
- Filtering